

# C++ London

# University Session 22

Tristan Brindle

# Feedback and Communication

- Your feedback is vital
- Otherwise, we don't know what you don't know!
- Please join the #ug\_uk\_cpplondonuni channel on the cpplang Slack — Go to <https://cpplang.now.sh/> for an “invitation”

# Today's Lesson Plan

- Introduction to the STL, part 2
  - Iterators revision
  - Algorithmic complexity
  - Overview of STL containers
  - Overview of STL algorithms
- Exercise

# STL Overview

- The **standard template library** (STL) was invented by Alexander Stepanov in the early 1990s
- It provided a set of **container classes** and fundamental **algorithms**
- The STL pioneered the concept of *generic programming*, revolutionising the way C++ was written and used
- Stepanov's STL formed the basis for much of the C++ standard library that we use today

# Revision: Iterators

- Iterators are the “glue” that binds together STL containers and algorithms
- Algorithms don’t operate on containers directly; rather, they operate on iterators which containers produce
- An iterator represents a **position** in a sequence.
- We can *dereference* an iterator to access the element at its position, and *advance* the iterator so that it moves to the next position

# Revision: Iterators

- Algorithms operate on iterator pairs: the first iterator points to the start of the sequence, and second points to one place **past the end** of the sequence
- (In mathematical terms, an iterator pair is a *half-open range*)
- There is work underway to add range-based algorithms to the standard library, so you can just say

```
std::vector<int> vec{5, 4, 3, 2, 1};  
std::ranges::sort(vec);
```

# Revision: Iterators

- There are five *categories* of iterator, forming a hierarchy:
  - Input/output: single-pass, deref, advance, compare
  - Forward: multi-pass
  - Bidirectional: can step backwards
  - Random access: can step forwards or backwards arbitrary distances
- Some algorithms require a certain category of iterator (e.g. `std::sort()` requires random access iterators)
- Other algorithms provide a more efficient implementation if passed a higher iterator category

# Algorithmic Complexity

- In computer science, the *complexity* of an algorithm refers to the amount of resources required to execute it
- We are mostly interested in *time complexity*, though *space complexity* is also important for some problems
- Algorithmic analysis is a complicated and deeply theoretical subject. We will only touch on the very very basics today
- For a more rigorous mathematical treatment, consult a computer science textbook 😊



# Big-O Notation

- Most of the time, we are interested in how an algorithm's performance scales as we increase the problem size
- We can express this using “big-O” notation
- Roughly, given a (mathematical) function  $f(n)$ , its behaviour as  $n$  increases is dominated by the fastest-growing term
- We call this the *order* of the function, denoted  $O(x)$

# Big-O Notation

- For example, if a given algorithm on  $n$  elements takes


$$f(n) = 4n^2 + 3n + 8$$

operations, then as  $n$  grows very large  $f(n)$  will be dominated by the  $n^2$  term

- We say that this algorithm has order  $n^2$ , or  $O(n^2)$

# Big-O Notation

- If an algorithm does not depend on the number of elements, we say it operates in *constant time*, written  $O(1)$
- An  $O(n)$  algorithm is said to operate in linear time
- Other common complexity classes:
  - $O(\log n)$ : logarithmic time
  - $O(n^2)$ : quadratic time
  - $O(n^k)$ : polynomial time
  - $O(2^n)$ : exponential time



By Cmglee - Own work, CC BY-SA 4.0, <https://commons.wikimedia.org/w/index.php?curid=50321072>

# Big-O is not everything!

- The theoretical complexity is important in choosing an appropriate algorithm
- However, “big-O” masks constant factors and lower-order terms, which can be more important for real-world problem sizes
- In particular, caching and modern CPU optimisations can have surprising effects
- When in doubt — **measure!**

# Complexity and the STL

- For containers and algorithms, the C++ standard does not specify exactly the implementation that should be used
- Rather, it specifies the **algorithmic complexity** that operations must have
- For example, `std::vector`'s operator[] must operate in constant time. `std::sort()` must have  $O(n \log n)$  complexity.

# Standard Containers

- The containers in the standard library can be broadly divided into four categories:
  - Sequence containers: elements are stored in the order they are added
  - Associative containers: elements are sorted for fast lookup
  - Unordered associative containers: elements are hashed for fast lookup
  - Container adaptors: provide a different interface for specific tasks

# Standard Containers

- All standard containers are class templates which meet the standard library's Container concept.
- In particular, this means they had nested types named `iterator` and `const_iterator` which meet the Iterator requirements
- Note that `iterator` and `const_iterator` need not be the same (and usually are not). However, you can always convert an `iterator` into a `const_iterator`.
- Calling `begin()` or `end()` on a non-const instance of a container will return an `iterator`. Calling `begin()` or `end()` on a const instance of a container will return a `const_iterator`.
- All standard containers take an `Allocator` template parameter: this can be used to optimise memory allocation for advanced uses.



# Sequence Containers

- `std::array`: Fixed-size random-access array
- `std::vector`: Dynamically sized random-access array
- `std::list`: Doubly-linked list
- `std::forward_list`: Singly-linked list
- `std::deque`: Double-ended queue
- `std::string`: Dynamically-sized random-access array

# std::vector

- `std::vector` is the standard's version of a dynamically-sized, random-access array
- Properties:
  - Element access is constant time
  - Adding elements to the end of a vector is *amortised constant time*
  - Adding elements anywhere else is *linear* in the number of elements in the vector
  - Iterators are random-access (contiguous in C++17)
- `std::vector` should be your default, go-to container for almost everything
- **When in doubt, use vector!**
- Further: if you can't use vector, think about how to change your data structures so that you can

# std::list

- `std::list` is a doubly-linked list
- Properties:
  - Adding/removing an element anywhere in the list is constant time
  - Accessing a particular element is linear in the number of elements in the list
  - Accessing the `size()` of the list is constant time in C++11
  - Iterators are bidirectional
- Note that `std::list` is *node-based*, meaning adding an element requires a dynamic allocation
- Guideline: prefer `std::list` to `std::vector` only when frequently inserting and removing elements from the middle of a *large* sequence
- Even then, try profiling first!

# std::deque

- `std::deque` is a double-ended queue
- Properties:
  - Element access is constant time
  - Insertion or removal at the end **or the beginning** of a deque is constant time
  - Insertion or removal elsewhere is linear in the size of the deque
  - Iterators are random access
- `std::deque` offers the same complexity guarantees as `std::vector`, plus an extra guarantee regarding insertion at the start
- But there's no such thing as a free lunch: deque's operations have a larger constant factor, and there is likely to be extra memory overhead

# Associative Containers

- `std::map`: A collection of unique key-value pairs, sorted by keys
- `std::set`: A sorted collection of unique keys
- `std::multimap`: A collection of key-value pairs, sorted by keys
- `std::multiset`: A sorted collection of keys

# `std::map`

- `std::map` is a sorted container of unique key-value pairs
- Properties:
  - Element access and insertion is  $O(\log n)$
  - Iterators are bidirectional
  - Iteration order is well-defined
- Typically implemented as a red-black tree
- By default, keys are sorted using operator<. This can be customised using map's Compare template parameter
- `std::map` should be your default, go-to associative array type for most purposes

# `std::set`

- A `std::set` is a sorted container of unique values
- Think of it as a map without the values!
- You can often achieve better performance by using a `std::vector` and keeping it sorted yourself

# Unordered Associative Containers

- `std::unordered_map`
- `std::unordered_set`
- `std::unordered_multimap`
- `std::unordered_multiset`



# `std::unordered_map`

- `std::unordered_map` is the standard library's version of a hash table
- Properties:
  - Element access and insertion:  $O(1)$  average,  $O(n)$  worst-case
  - Iterators are forward only
  - Iteration order is undefined
- Compared with `std::map`, `unordered_map` offers constant-time lookup on average
- Keys are compared using `std::hash` by default. This is defined for all built-in types, but you need to specialise it for your own types
- As with all hash tables, performance is highly dependent on the quality of the hash function

# Container adaptors

- The container adaptors in the STL take an underlying sequence container as a template argument and wrap it in a new, more specialised, more restrictive API
- The adapted classes are not containers themselves (they have no `begin()` or `end()`), so in practise are rarely used
- The container adaptors are:
  - `std::queue`: Adapts a container (such as vector) into a FIFO queue
  - `std::stack`: Adapts a container into a LIFO stack
  - `std::priority_queue`: Adapts a container so as to provide constant-time access to the largest element

# Standard Algorithms

- The standard library contains over 80 different algorithms!
- These are mostly defined in header `<algorithm>`, with some extras like `std::accumulate()` defined in header `<numeric>`
- These range from the very basic (e.g. `std::count()`) to the very specialised (e.g. `std::sort()`)
- General advice: use the standard algorithms whenever you can!
- Further: if you find yourself writing a non-trivial for loop, see if you can abstract out the operation into a self-contained algorithm
- See Sean Parent's fantastic C++ Seasoning talk for inspiration

# Predicates

- Many standard algorithms allow you to pass *predicates* which dictate their behaviour
- A predicate is just a callable (a function, function object or lambda) which returns a bool
- Some algorithms require unary predicates taking one argument (usually of the iterator's `value_type`)
- Other algorithms require a binary predicate, comparing two elements of the range

# Using Predicates

```
struct Person {
    std::string first_name;
    std::string last_name;
};

bool compare_person(const Person& a, const Person& b)
{
    if (a.last_name == b.last_name) {
        return a.first_name < b.first_name;
    }
    return a.last_name < b.last_name;
}

std::vector<Person> people;
/* ...fill people vector... */

// Get iterator to the first person, sorted alphabetically
auto iter = std::min_element(people.begin(), people.end(), compare_person);

// Dereference the iterator to get a reference to the Person instance
Person& first_person = *iter;
```

# Anatomy of a standard algorithm

```
template <typename Iter, typename UnaryPredicate>
std::size_t count_if(Iter first, Iter last, UnaryPredicate pred)
{
    std::size_t count = 0;

    while (first != last) {
        if (pred(*first)) {
            ++count;
        }
        ++first;
    }

    return count;
}

template <typename Iter, typename T>
std::size_t count(Iter first, Iter last, const T& value)
{
    using value_type = typename std::iterator_traits<Iter>::value_type;
    return count_if(first, last, [&value] (const value_type& i) {
        return i == value;
    });
}
```

# Exercise

- [https://github.com/CPPLondonUni/algorithms\\_exercise](https://github.com/CPPLondonUni/algorithms_exercise)

# Online Resources

- <https://isocpp.org/get-started>
- [cppreference.com](http://cppreference.com) — The bible, but aimed at experts
- [cplusplus.com](http://cplusplus.com) — Another reference site, also has a tutorial section
- [learncpp.com](http://learncpp.com) — Free online tutorial, very up-to-date
- <https://www.pluralsight.com/authors/kate-gregory> - Comprehensive set of courses from an experienced C++ trainer (free trial)
- [reddit.com/r/cpp\\_questions](https://reddit.com/r/cpp_questions)
- Cpplang Slack channel — <https://cpplang.now.sh/> for an “invite”
- StackOverflow (but...)



# Thanks for coming!

C++ London University:

- Website: [cpplondonuni.com](http://cpplondonuni.com)
- Twitter: @cpplondonuni
- Github: [github.com/CPPLondonUni](https://github.com/CPPLondonUni)
- Reddit: [reddit.com/r/CppLondonUni](https://reddit.com/r/CppLondonUni)

Where to find Tom Breza:

- On Slack: [#learn #ug\\_uk\\_cpplondonuni](https://cpplang.slack.com)
- E-mail: [tom@PCServiceGroup.co.uk](mailto:tom@PCServiceGroup.co.uk)
- Mobile: 07947451167

My stuff:

- Website: [tristanbrindle.com](http://tristanbrindle.com)
- Twitter: @tristanbrindle
- Github: [github.com/tcbrindle](https://github.com/tcbrindle)

See you next time! 😊